

Foveated Memory: Constant-Cost Attention with Sinks and a Sliding Window

20 Watts · Episode 3 — Foveated Memory

[Your Name], age 17 — June 2026 · Code & data: github.com/svaka2000/20-watts

Abstract

Human memory does not retain a verbatim transcript of everything we perceive; it keeps the gist and the most recent detail, and lets the rest fade. A transformer does the opposite: it stores a perfect key/value (KV) cache for **every** past token, so both memory and attention cost grow without bound with context length. We test a brain-like alternative on **Qwen2.5-7B-Instruct (4-bit)**: keep only a few initial **attention-sink** tokens plus a sliding window of the most recent W tokens, and drop the rest (the StreamingLLM recipe; Xiao et al., 2023). Implemented as a custom attention mask whose large-window limit is *bit-exact* with full causal attention (faithfulness check: $|\Delta ppl| = 0.00$), we measure held-out perplexity vs the retained KV budget. Keeping just **26% of the cache ($W=512$) costs +34% perplexity**; keeping 7% costs +71%. The cost is real and grows as you forget more — an honest result, not a free lunch. The striking finding is the **attention sink**: at $W=128$, removing the first few tokens explodes perplexity to **+542%**, while restoring a **single** sink token collapses it back to **+73%**. The value of foveated memory is therefore not "free compression" but **bounded, constant memory for unbounded context** — something full attention cannot offer at any price.

1. Introduction

Attention is the part of a transformer whose cost scales with how much you remember. For a context of T tokens, the KV cache holds $O(T)$ state per layer and each new token attends over all T . Brains face the same pressure and solve it by *forgetting selectively*: the hippocampus preserves salient anchors and the gist, not a literal recording. This episode asks how much a language model can forget, and how gracefully.

We adopt the **StreamingLLM** memory (Xiao et al., 2023): a fixed budget of S "sink" tokens at the very start plus the most recent W tokens. Contribution: a faithful, reproducible measurement of the resulting quality/memory trade-off on a modern 7B model, including a clean reproduction of the **attention-sink phenomenon**.

2. Method

We replace the model's causal attention mask with a streaming mask: query position i may attend to key j iff $j \leq i$ and $(j < S \text{ or } j > i - W)$. All other entries get a large negative bias. This is applied globally to every layer by overriding the model's mask constructor; nothing else changes.


Faithfulness check. With $W \rightarrow \infty$ the streaming mask is exactly causal, so perplexity must equal the unmodified model. Measured: causal ppl = 5.914, streaming($W=\infty$) ppl = 5.914, $|\Delta| = 0.00$ — the harness is exact, so any change at finite W is the genuine cost of forgetting.

Evaluation: WikiText-2 test, 1951 contiguous tokens.

3. Results

3.1 Memory vs quality

Window W	KV kept (of 1951)	Memory saved	Perplexity	Δ vs full
∞ (full)	1951	0%	5.914	—
512	516	74%	7.943	+34.3%
256	260	87%	8.792	+48.7%
128	132	93%	10.119	+71.1%
64	68	97%	11.831	+100%
32	36	98%	15.735	+166%



On a coherent document, far-back context genuinely carries information, so forgetting it has a real, monotonic cost. We report this plainly: foveated memory is a *memory-bounded approximation*, not a free saving. Its purpose is the regime full attention cannot reach — streams longer than you can afford to store.

3.2 The attention sink (the interesting part)

Holding the window at $W=128$ and varying the number of sink tokens S :

Sinks S	Perplexity	Δ vs full
0	37.97	+542%
1	10.22	+72.8%
2	10.18	+72.1%
4	10.12	+71.1%

Dropping the first few tokens is *catastrophic*, and a **single** sink token recovers almost all of the loss. This reproduces Xiao et al.'s discovery: a trained transformer offloads a large amount of "default" attention onto the first positions; those tokens act as a pressure-release valve, and evicting them breaks the softmax. It is a vivid example of a model-internal mechanism you only see when you intervene — and our bit-exact harness makes the reproduction trustworthy.

4. Discussion

What stacks, what doesn't. Foveated memory targets *attention/memory* cost and is orthogonal to sparse firing (Episode 1, MLP compute) and quantization (storage). It composes with them (see the synthesis, `stack_all.py`). Its honest value: **constant memory and constant per-token attention cost at arbitrary context length**, plus a clean window to study attention sinks.

Limitations. Pure recency+sink eviction ignores *which* old tokens mattered; score-based methods (H2O, Zhang et al. 2023) keep "heavy hitter" tokens and trade off better. Our perplexity is on one corpus; long-stream stability (StreamingLLM's headline) is a separate, complementary metric.

5. Conclusion

A model can run on a *fixed* memory budget regardless of context length, at a quality cost that is real but bounded — and the humble first token turns out to be doing heroic work. Foveated memory is Episode 3 of *20 Watts*; combined with sparse firing and adaptive depth it forms the full stack toward a 20-watt machine.

References

1. G. Xiao, Y. Tian, B. Chen, S. Han, M. Lewis. *Efficient Streaming Language Models with Attention Sinks (StreamingLLM)*. arXiv:2309.17453, 2023.
2. Z. Zhang, Y. Sheng, T. Zhou, et al. *H2O: Heavy-Hitter Oracle for Efficient Generative Inference of LLMs*. NeurIPS, 2023.
3. P. Lennie. *The Cost of Cortical Computation*. Current Biology, 13(6):493–497, 2003.